

```
In [1]: from sklearn.datasets import load_iris
import pandas as pd
iris = load_iris()
df = pd.DataFrame(data=iris.data, columns=iris.feature_names)
df['species'] = iris.target
```

```
In [3]: print(f'Dataset Dimensions: {df.shape}')
```

Dataset Dimensions: (150, 5)

```
In [5]: print(df.head())
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
\				
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
2	4.7	3.2	1.3	0.2
3	4.6	3.1	1.5	0.2
4	5.0	3.6	1.4	0.2

	species
0	0
1	0
2	0
3	0
4	0

```
In [7]: print(df.describe())
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	\
count	150.000000	150.000000	150.000000	
mean	5.843333	3.057333	3.758000	
std	0.828066	0.435866	1.765298	
min	4.300000	2.000000	1.000000	
25%	5.100000	2.800000	1.600000	
50%	5.800000	3.000000	4.350000	
75%	6.400000	3.300000	5.100000	
max	7.900000	4.400000	6.900000	

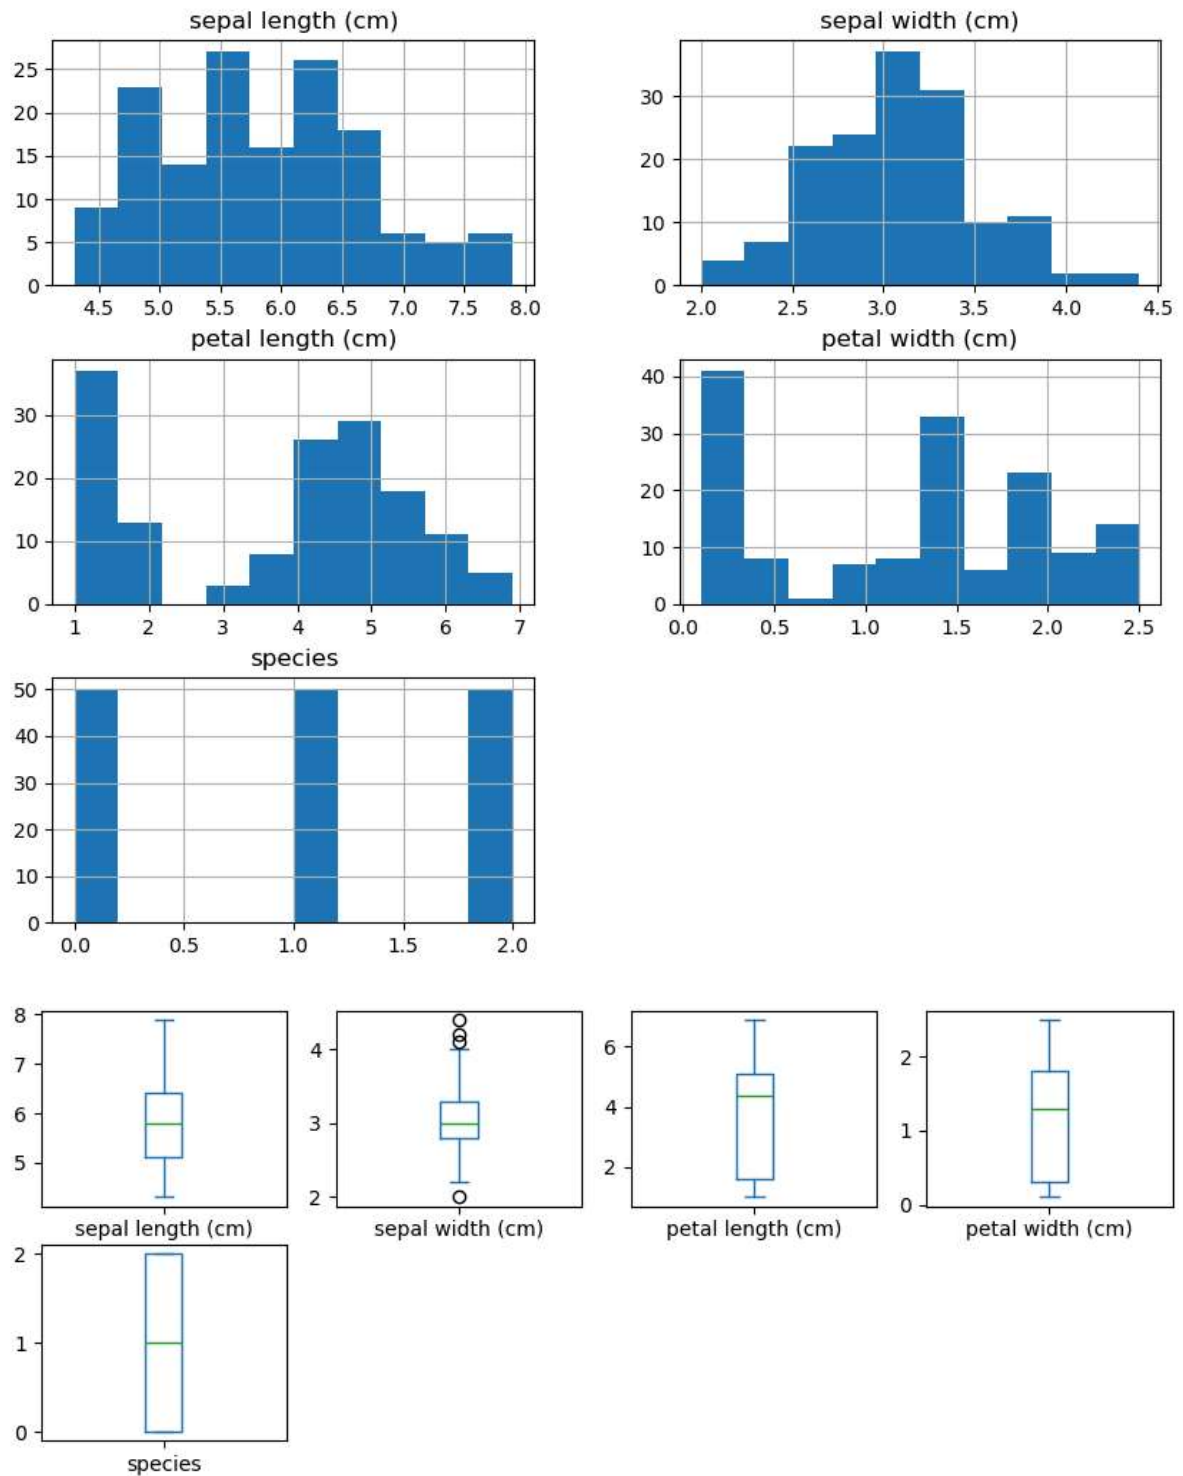
  

	petal width (cm)	species
count	150.000000	150.000000
mean	1.199333	1.000000
std	0.762238	0.819232
min	0.100000	0.000000
25%	0.300000	0.000000
50%	1.300000	1.000000
75%	1.800000	2.000000
max	2.500000	2.000000

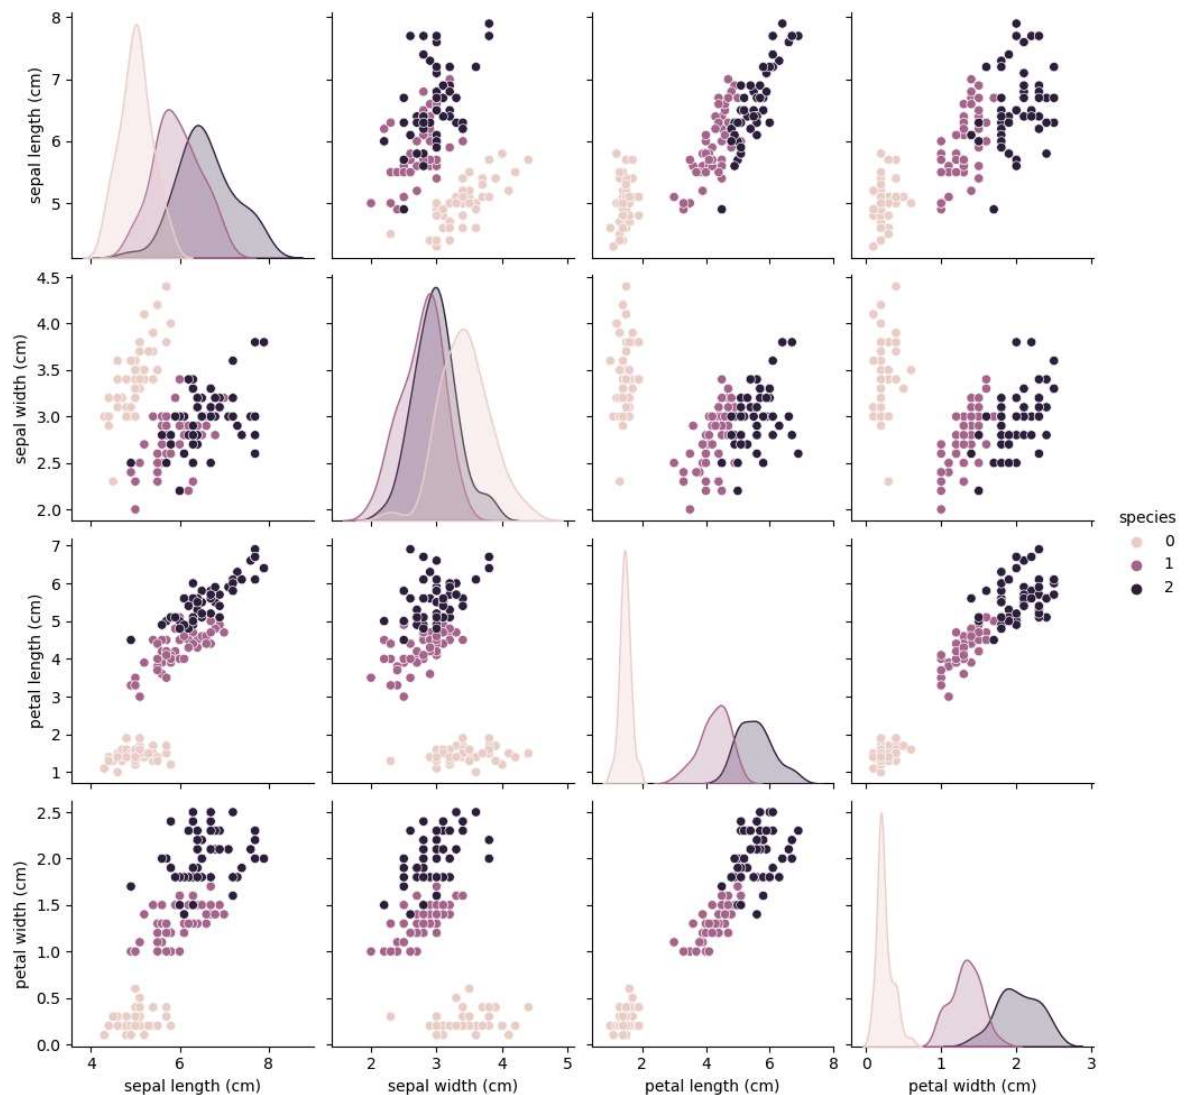
```
In [9]: print(df['species'].value_counts())
```


```
species
0      50
1      50
2      50
Name: count, dtype: int64
```

```
In [13]: import matplotlib.pyplot as plt
import seaborn as sns
df.hist(figsize=(10, 8))
plt.show()
df.plot(kind='box', subplots=True, layout=(4,4), figsize=(10, 8))
plt.show()
```



```
In [33]: import warnings
warnings.filterwarnings('ignore')
import seaborn as sns
import matplotlib.pyplot as plt
sns.pairplot(df, hue='species')
plt.show()
plt.figure(figsize=(40, 20))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.show()
from IPython.core.display import display, HTML
display(HTML("<style>.container { width:100% !important; }</style>"))
```





```
In [17]: from sklearn.model_selection import train_test_split
X = df.iloc[:, :-1]
y = df['species']
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_
```

```
In [29]: from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
models = []
models.append(('LR', LogisticRegression(max_iter=200)))
models.append(('LDA', LinearDiscriminantAnalysis()))
models.append(('KNN', KNeighborsClassifier()))
models.append(('CART', DecisionTreeClassifier()))
models.append(('NB', GaussianNB()))
models.append(('SVM', SVC()))
results = []
names = []
for name, model in models:
    cv_results = cross_val_score(model, X_train, y_train, cv=10, scoring='accu
    results.append(cv_results)
    names.append(name)
    print(f'{name}: {cv_results.mean()} ({cv_results.std()})')
```

LR: 0.9666666666666666 (0.04082482904638632)

LDA: 0.9749999999999999 (0.03818813079129868)

KNN: 0.9583333333333333 (0.041666666666666685)

CART: 0.9583333333333333 (0.041666666666666685)

NB: 0.9499999999999998 (0.05527707983925667)

SVM: 0.975 (0.03818813079129868)

```
In [31]: model = KNeighborsClassifier()
model.fit(X_train, y_train)
predictions = model.predict(X_val)
from sklearn.metrics import accuracy_score, classification_report
print(f'Accuracy: {accuracy_score(y_val, predictions)}')
print(classification_report(y_val, predictions))
```

Accuracy: 1.0

	precision	recall	f1-score	support
0	1.00	1.00	1.00	11
1	1.00	1.00	1.00	13
2	1.00	1.00	1.00	6
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30